

# Computer Networks - *Xarxes de Computadors*

## Outline

- Course Syllabus
- Unit 1: Introduction
- Unit 2. IP Networks
- **Unit 3. Point to Point Protocols -TCP**
- Unit 4. LANs
- Unit 5. Data Transmission

# Unit 3. Point to Point Protocols -TCP

## Outline

- **Introduction**
- Basic ARQ Protocols
- UDP Protocol
- TCP Protocol

## Unit 3. Point to Point Protocols -TCP

### Introduction

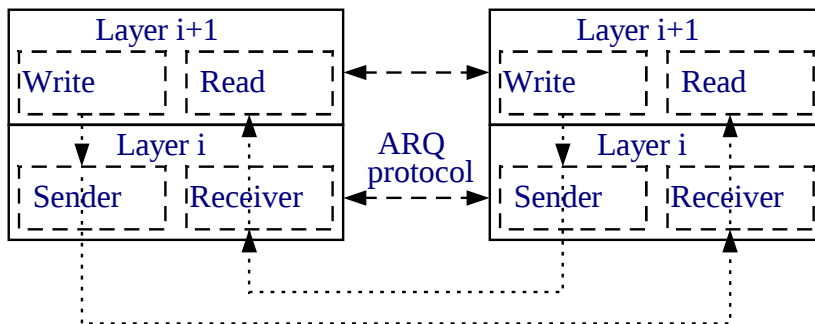
- A **Point to Point Protocol** (PPP) takes place between exactly two endpoints.
- PPP is usually used to identify protocols that builds up a communication channel between endpoints, adding **features** of the type:
  - **Error detection**
  - **Error recovery**
  - **Flow control**
- These are typical data-link layer features, although protocols from other **layers** can be also regarded as PPPs:
  - Physical: RS-232
  - Data-link: The PPP protocol used in TCP/IP
  - Network: X.25
  - Transport: TCP
- **Automatic Repeat reQuest (ARQ)** protocols are typically used for PPP.

## Unit 3. Point to Point Protocols -TCP


### Introduction - Automatic Repeat reQuest, ARQ

#### ARQ Ingredients

- Connection oriented
- Tx/Rx buffers
- Acknowledgments (ack)
- Acks can be *piggybacked* in information PDUs sent in the opposite direction.
- Retransmission Timeout.
- Sequence Numbers



ARQ Protocol Architecture



ARQ Protocol Implementation (one way)

## Unit 3. Point to Point Protocols -TCP

### Introduction - Automatic Repeat reQuest, ARQ

#### Basic ARQ Protocols:

- Stop & Wait
- Go Back N
- Selective Retransmission

# Unit 3. Point to Point Protocols -TCP


## Outline

- Introduction
- **Basic ARQ Protocols**
- UDP Protocol
- TCP Protocol

## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols - Assumptions

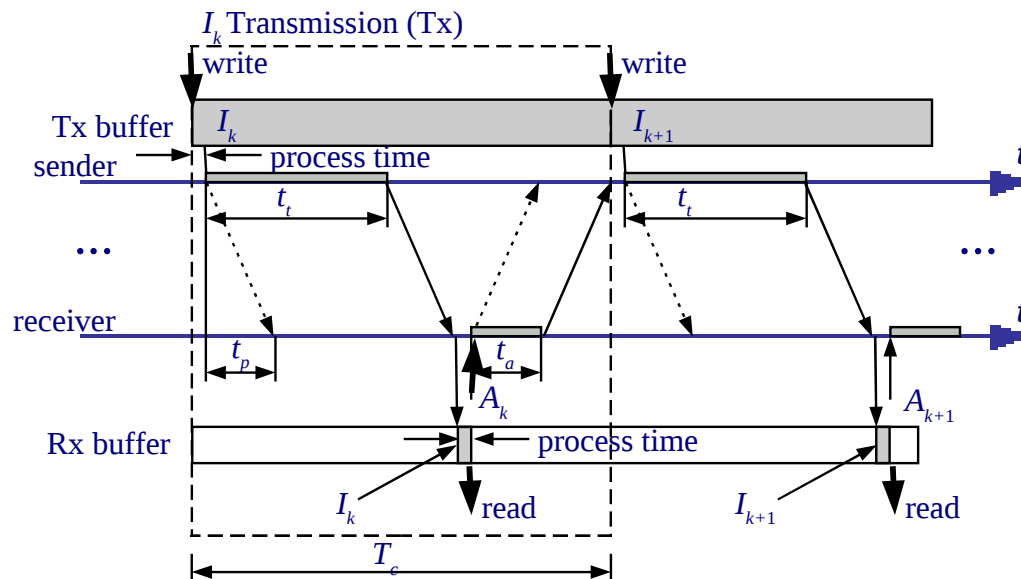
- We shall focus on the the transmission in **one direction**.
- We shall assume a **saturated source**: There is always information ready to send.
- We shall assume **full duplex** links.
- ppp protocol over a **line of  $D$  m distance and  $v_t$  bps bitrate**.
- **Propagation speed of  $v_p$  m/s**, thus, propagation delay of  $D/v_p$  s.
- We shall refer to a generic layer, where the sender sends **Information PDUs ( $I_k$ )** and the receiver sends **ack PDUs ( $A_k$ )**.
- Frames carrying  $I_k$  respectively  $A_k$ , are Tx using  $L_I$  and  $L_A$  bits, thus the **Tx times** are respectively:  $t_t = L_I/v_t$  and  $t_a = L_A/v_t$  s.




## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols - Stop & Wait

1. **When the sender is ready:** (i) allows writing from upper layer, (ii) builds  $I_k$ , (iii)  $I_k$  goes down to data-link layer and Tx starts.
2. **When  $I_k$  completely arrives to the receiver:** (i) it is read by the upper layer, (ii)  $A_k$  is generated,  $A_k$  goes down to data-link layer and Tx starts.
3. When  $A_k$  completely arrives to the sender, goto 1.



Time diagram




Simplified time diagram



## Unit 3. Point to Point Protocols -TCP


### Basic ARQ Protocols - Stop & Wait Retransmission

- Each time the sender Tx a PDU, a **retransmission timeout (TO)** is started.
- If the information PDU do not arrives, or arrives with **errors**, **no ack is sent**.
- When TO expires, the sender **ReTx the PDU**.




## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Why sequence numbers are needed?



Need to number information PDUs




Need to number ack PDUs

## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Notes on computing the efficiency (channel utilization)

- Line bitrate:  $v_t = 1/t_b$ , bps
- **Throughput** (*velocidad efectiva*)  $v_{ef}$  = number of inf. bits / obs. time, bps
- **Efficiency** or channel utilization  $E = v_{ef} / v_t$  (times 100, in percentage)




$$E = \frac{v_{ef}}{v_t} = \frac{\# \text{info bits} / T}{1/t_b} = \left\{ \begin{array}{l} \frac{\# \text{info bits} \times t_b}{T} = \frac{\text{time Tx information}}{T} \\ \frac{\# \text{info bits}}{T/t_b} = \frac{\# \text{info bits}}{\# \text{bits at line bitrate}} \end{array} \right.$$

## Unit 3. Point to Point Protocols -TCP


### Basic ARQ Protocols – Stop & Wait efficiency

- Assuming no errors (maximum efficiency), the Tx is periodic, with period  $T_c$ .
- $E_{protocol}$ : We do not take into account headers.



$$E_{protocol} = \frac{t_t}{T_c} = \frac{t_t}{t_t + t_a + 2t_p} =$$


$$\frac{t_t}{t_t + 2t_p} \simeq \frac{1}{1 + 2a}, \text{ where } a = \frac{t_p}{t_t}$$



## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Continuous Tx Protocols


- Goal: Allow high efficiency independently of propagation delay.
- Without errors:  $E = 100\%$



## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Go Back N


- **Cumulative acks:**  $A_k$  confirm  $I_i$ ,  $i \leq k$
- If the sender receives an error of out of order PDU: Do not send acks, discards all PDU until the expected PDU arrives. Thus, **the receiver does not store out of order PDUs.**
- When a **TO** occurs, the sender *goes back* and starts Tx from that PDU.



## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Selective ReTx.


- The same as Go Back N, but:
  - The sender only ReTx a PDU when a TO occurs.
  - The **receiver stores out of order PDUs**, and ack all stored PDUs when missing PDUs arrive.



## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Efficiency with Tx errors: Stop & Wait

- Assumptions: On average,  $N_t$  Tx are needed to successfully send a PDU:  $N_t - 1$  with Tx errors, and 1 correct.



$$E_{protocol} = \frac{t_t}{T_T} = \frac{t_t}{(N_t - 1)T_O + T_C}$$

- To avoid unnecessary ReTx  $T_O > T_C$ . The maximum throughput is when  $T_O \approx T_C$ :


$$E_{protocol} \simeq \frac{t_t}{N_t T_C} = \frac{1}{N_t (1 + 2a)}$$



## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Efficiency with Tx errors: Go Back N

- Assumptions: On average,  $N_t$  Tx are needed to successfully send a PDU:  $N_t - 1$  with Tx errors, and 1 correct.



$$E_{protocol} = \frac{t_t}{T_T} = \frac{t_t}{(N_t - 1)T_O + t_t}$$


- The maximum throughput is when  $T_O \approx T_C$ :

$$E_{protocol} \simeq \frac{t_t}{(N_t - 1)T_C + t_t} = \frac{1}{(N_t - 1)(1 + 2a) + 1} = \frac{1}{N_t(1 + 2a) - 2a}$$

## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Efficiency with Tx errors: Select. ReTx

- Assumptions: On average,  $N_t$  Tx are needed to successfully send a PDU:  
 $N_t - 1$  with Tx errors, and 1 correct.




$$E_{\text{protocol}} = \frac{t_t}{N_t t_t} = \frac{1}{N_t}$$

## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Efficiency with Tx errors: Comparison


- For  $a \ll 1$ , all protocols are similar ( $a = t_p/t_t$ )
- For  $N_t \gg 1$  all protocols have low  $E$ .
- $E$  in selective ReTx does not depend on  $a$ .
- If  $a \ll 1$  does not hold:
  - $E$  of Stop & Wait is low.
  - For moderate  $N_t$ , Sel-ReTx outperforms GoBackN.



## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Flow Control and Window Protocols

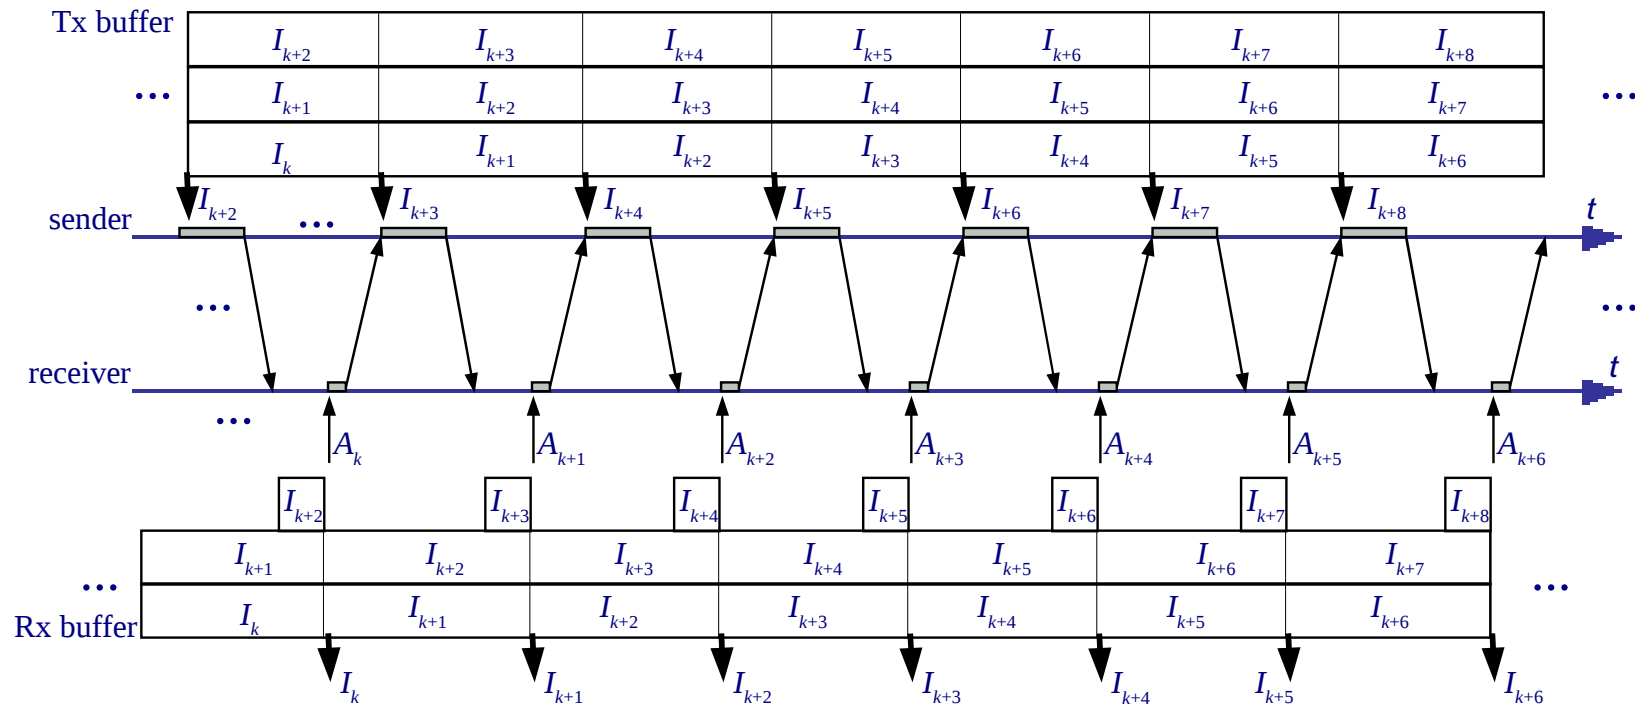
- ARQ are also used for flow control. **Flow control** consists on avoiding the sender to Tx at higher PDU rate than can be consumed by the receiver.
- With **Stop & Wait**, if the receiver is slower, acks are delayed and the sender reduces the throughput.
- With continuous Tx protocols: A **Tx window** is used. The window is the maximum number of non-ack PDUs that can be Tx. If the Tx window is exhausted, the sender stales.
- **Stop & Wait** is a window protocol with **Tx window = 1 PDU**.
- Furthermore, the Tx window allows **dimensioning the Tx buffer**, and the Rx buffer for Selective ReTx: No more the Tx window PDUs need to be stored.



# Unit 3. Point to Point Protocols -TCP

## Basic ARQ Protocols – Window Protocol Flux Control Example


- $W = 3$
- Stationary regime




## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Optimal Tx window

- **Optimal window:** Minimum window that allows the maximum throughput.
- Optimal window example:



- Non optimal window example:



- Clearly, for this example:

$$W_{opt} = \left\lceil \frac{T_c}{t_t} \right\rceil$$

## Unit 3. Point to Point Protocols -TCP

### Basic ARQ Protocols – Sequence Number Field Dimensioning

- Information and ack PDUs need **sequence numbers**.
- Using  $n$  bits, the sequence number used for PDU  $k$  is:  $k \bmod 2^n$ .
- What value we need for  $n$  to avoid ambiguity?
- Without proof: If PDUs arrive in the same order they are Tx, with arbitrary delays, we need to distinguish the PDU that can be in the Tx and Rx buffers of the sender and receiver respectively:

Protocol	#Seq. Numbers ( $\geq$ )	#Bits ( $\geq$ )
Stop & Wait	2	1
GoBackN	$W+1$	$\log_2(W+1)$
Selective ReTx	$2 W$	$\log_2(2 W)$

# Unit 3. Point to Point Protocols -TCP

## Outline


- Introduction
- Basic ARQ Protocols
- **UDP Protocol**
- TCP Protocol



## Unit 3. Point to Point Protocols -TCP

### UDP Protocol – Introduction: The Internet Transport Layer

- Two protocols are used at the TCP/IP transport layer: User Datagram Protocol (**UDP**) and Transmission Control Protocol (**TCP**).
- **UDP** offers a *datagram service* (non reliable).
- **TCP** offers a reliable service.
- Transport layer offers a communication channel between applications.
- Transport layer access points (applications) are identified by a 16 bits port numbers.
- TCP/UDP use the client/server paradigm



## Unit 3. Point to Point Protocols -TCP

### UDP Protocol – Description (RFC 768)

- **UDP Datagram service:** same as IP.
  - Non reliable
  - No error recovery
  - No ack
  - Connectionless
  - No flux control
- UDP PDU is referred to as **UDP datagram**.
- UDP does not have a Tx buffer: each application write operation generates a UDP datagram.
- UDP is typically used:
  - Applications where short messages are exchanged: E.g. DHCP, DNS, RIP.
  - Real time applications: E.g. Voice over IP, videoconferencing, stream audio/video. These applications does not tolerate large delay variations (which would occur using an ARQ).

# Unit 3. Point to Point Protocols -TCP

## UDP Protocol – C Code example

```
main(int argc, char *argv[])
{
    int sock ;
    struct sockaddr_in clnt_addr, serv_addr ;
    struct hostent *host ;
    int serv_len = sizeof(serv_addr) ;

    if(argc != 2) {
        fprintf(stderr, "usage: %s hostname\n", argv[0]) ;
        exit(1) ;
    }
    host = gethostbyname(argv[1]) ; /* call the resolver for server addr. */
    if(host == NULL) {
        perror("gethostbyname ") ;
        exit(2) ;
    }
    serv_addr.sin_family = AF_INET ;
    memcpy(&serv_addr.sin_addr, host->h_addr, host->h_length) ;
    serv_addr.sin_port = htons(3333) ;


    sock = socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP) ; /* create a socket */
    clnt_addr.sin_family = AF_INET ;
    clnt_addr.sin_addr.s_addr = htonl(INADDR_ANY) ; /* the OS choose any valid address */
    clnt_addr.sin_port = htons(0) ; /* the OS choose an ephemeral port */
    bind(sock, (struct sockaddr *)&clnt_addr, sizeof(clnt_addr)) ; /* give local addr. to a socket */

    char msg[] = "hello world\n" ;
    sendto(sock, msg, strlen(msg), 0,
           (struct sockaddr *)&serv_addr, sizeof(serv_addr)) ; /* send a UDP datagram */
    close(sock) ; /* close the socket */
}
```


## Unit 3. Point to Point Protocols -TCP

### UDP Protocol – UDP Header

- Fixed size of **8 bytes**.
- The **checksum** is computed using (i) the header, (ii) a pseudo-header, (iii) the payload.
- Because of the **pseudo-header**, the UDP checksum needs to be updated if NAT is used.



UDP datagram header



UDP pseudo-header

# Unit 3. Point to Point Protocols -TCP

## Outline

- Introduction
- Basic ARQ Protocols
- UDP Protocol
- **TCP Protocol**

## Unit 3. Point to Point Protocols -TCP

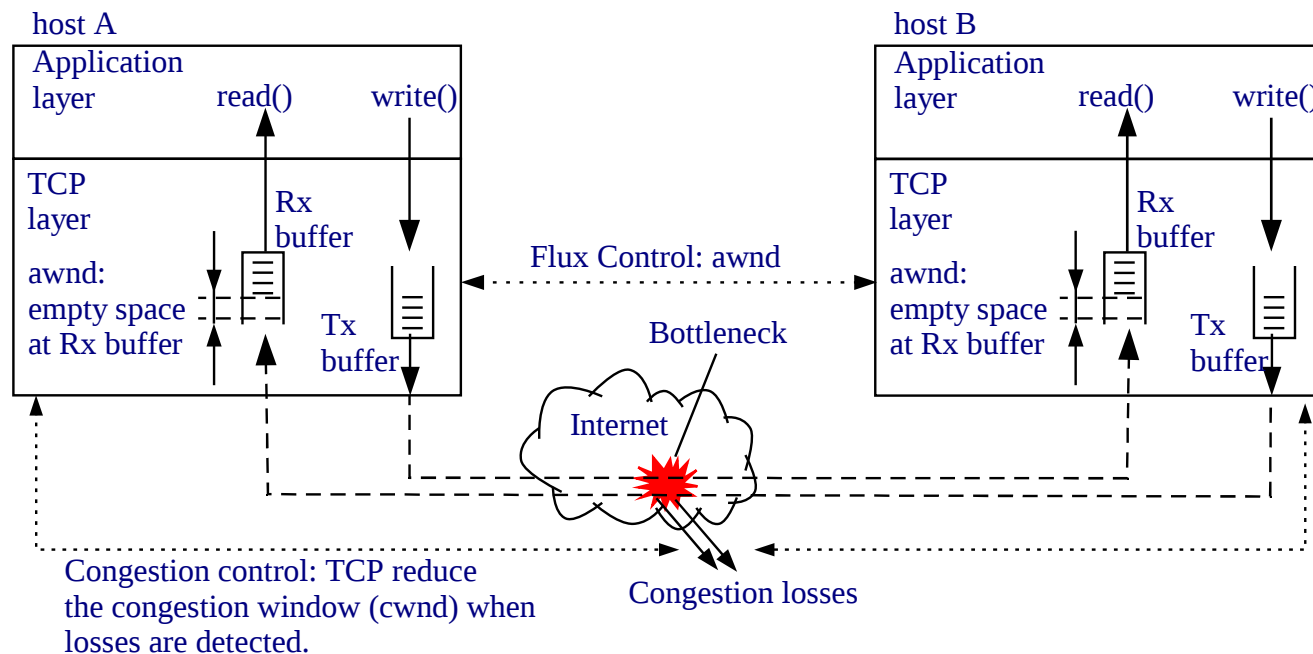
### TCP Protocol – Description (RFC 793)

- Reliable service (**ARQ**).
  - Error recovery
  - Acknowledgments
  - Connection oriented
  - Flux control
- TCP PDU is referred to as **TCP segment**.
- **Congestion control**: Adapt the TCP throughput to network conditions.
- Segments of optimal size: Variable Maximum Segment Size (**MSS**).
- TCP is typically used:
  - Applications requiring reliability: Web, ftp, ssh, telnet, mail, ...

## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Basic operation

- ARQ window protocol, with **variable window**:  $wnd = \min(awnd, cwnd)$
- Each time a segment arrives, TCP send an **ack** (unless delayed ack is used) without waiting for the upper layer to read the data.
- The advertised window (**awnd**) is used for flux control.
- The congestion window (**cwnd**) is used for congestion control.



## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Delayed acks and Nagle algorithm

- TCP connections can be classified as:
  - **Bulk:** (e.g. web, ftp) There are always bytes to send. TCP send MSS bytes.
  - **Interactive:** (eg. telnet, ssh) The user interacts with the remote host.
- In interactive connections small packets are sent: Each keyboard hit may generate a segment, and one ack is sent for each.
- Solutions: Delayed acks, Nagle algorithm.



## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Delayed acks and Nagle algorithm

- **Delayed ack.** It is used to reduce the amount of acks. Consists of sending **1 ack each 2 MSS segments, or 200 ms**. Acks are always sent in case of receiving out of order segments.

- tcpdump example:

bulk transfer



```

...
11:27:13.798849 147.83.32.14.ftp > 147.83.35.18.3020: P 9641:11089(1448) ack 1 win 10136 (DF)
11:27:13.800174 147.83.32.14.ftp > 147.83.35.18.3020: P 11089:12537(1448) ack 1 win 10136 (DF)
11:27:13.800191 147.83.35.18.3020 > 147.83.32.14.ftp: . 1:1(0) ack 12537 win 31856 (DF)
11:27:13.801405 147.83.32.14.ftp > 147.83.35.18.3020: P 12537:13985(1448) ack 1 win 10136 (DF)
11:27:13.802771 147.83.32.14.ftp > 147.83.35.18.3020: P 13985:15433(1448) ack 1 win 10136 (DF)
11:27:13.802788 147.83.35.18.3020 > 147.83.32.14.ftp: . 1:1(0) ack 15433 win 31856 (DF)
...

```

Diagram illustrating the structure of a TCP packet header and its flags:

```

timestamp      src IP addr/port      dst IP addr/port      TCP flags      seq. num:next seq      ack      awnd      DF flag in IP
                  |              |              |              |              |              |              |
11:27:13.798849  147.83.32.14.ftp > 147.83.35.18.3020: P 9641:11089(1448) ack 1 win 10136 (DF)

```


The diagram shows the following fields and their corresponding values in the packet header:

- timestamp:** 11:27:13.798849
- src IP addr/port:** 147.83.32.14.ftp
- dst IP addr/port:** 147.83.35.18.3020
- TCP flags:** P (Push flag)
- seq. num:next seq num (bytes):** 9641:11089(1448)
- ack:** 1
- awnd:** 10136
- DF flag in IP header set:** (DF)

## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Delayed acks and Nagle algorithm

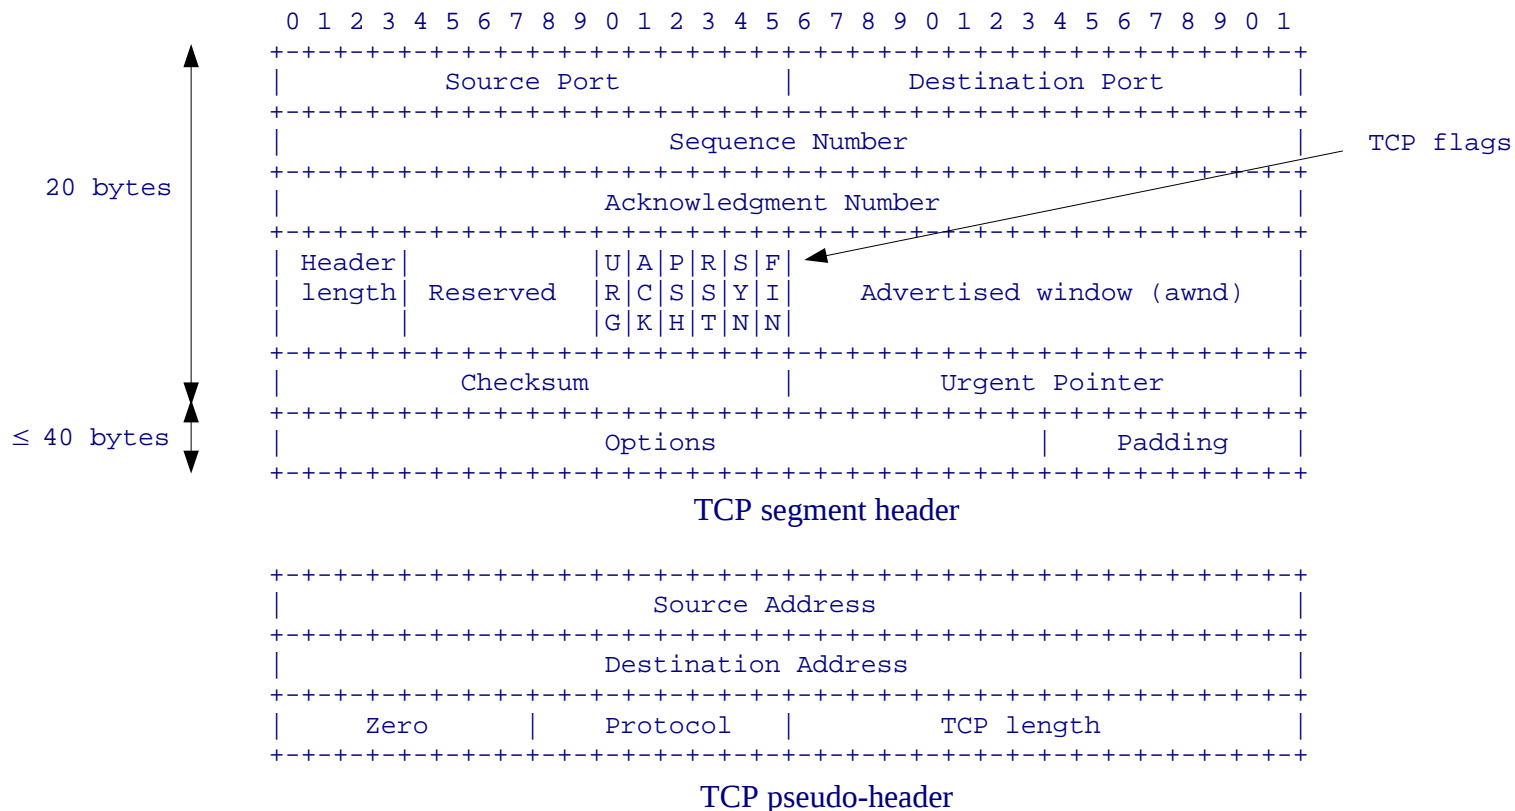
- **Nagle algorithm.** It is used to reduce the number of small information segments in iterative connections. Consists of TCP sending segments only when:
  - A full MSS segment can be sent.
  - There are no bytes pending to be ack. So, keyboard hits are stored until an ack arrives.



## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – TCP Header

- **Variable size:** Fixed fields of 20 bytes + options (15x4 = 60 bytes max.).
- Like UDP, the **checksum** is computed using (i) the header, (ii) a pseudo-header, (iii) the payload, and needs to be updated if PAT is used.



## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – TCP Flags

- **URG** (Urgent): The Urgent Pointer is used. It points to the first urgent byte. Rarely used. Example: ^C in a telnet session.
- **ACK**: The ack field is used. Always set except for the first segment sent by the client.
- **PSH** (Push): The sender indicates to “push” all buffered data to the receiving application. Most BSD derived TCPs set the PSH flag when the send buffer is emptied.
- **RST** (Reset): Abort the connection.
- **SYN**: Used in the connection setup (*three-way-handshaking, TWH*).
- **FIN**: Used in the connection termination.

# Unit 3. Point to Point Protocols -TCP

## TCP Protocol – TCP Flags

- tcpdump example:

TCP flags

S: SYN

P: PUSH

∴ No flag (except ack) is set

```

09:33:02.556785 IP 147.83.34.125.24374 > 147.83.194.21.80: S 3624662632:3624662632(0) win 5840
<mss 1460,sackOK,timestamp 531419155 0,nop,wscale 7>
09:33:02.558054 IP 147.83.194.21.80 > 147.83.34.125.24374: S 2204366975:2204366975(0) ack
3624662633 win 5792 <mss 1460,sackOK,timestamp 3872304344 531419155,nop,wscale 2>
09:33:02.558081 IP 147.83.34.125.24374 > 147.83.194.21.80: . ack 1 win 46 <nop,nop,timestamp
531419156 3872304344>
09:33:02.558437 IP 147.83.34.125.24374 > 147.83.194.21.80: P 1:627(626) ack 1 win 46
<nop,nop,timestamp 531419156 3872304344>
09:33:02.559146 IP 147.83.194.21.80 > 147.83.34.125.24374: . ack 627 win 1761 <nop,nop,timestamp
3872304345 531419156>
09:33:02.559507 IP 147.83.194.21.80 > 147.83.34.125.24374: P 1:271(270) ack 627 win 1761
<nop,nop,timestamp 3872304345 531419156>
09:33:02.559519 IP 147.83.34.125.24374 > 147.83.194.21.80: . ack 271 win 54 <nop,nop,timestamp
531419156 3872304345>
09:33:02.560154 IP 147.83.194.21.80 > 147.83.34.125.24374: . 271:1719(1448) ack 627 win 1761
<nop,nop,timestamp 3872304345 531419156>
09:33:02.560167 IP 147.83.34.125.24374 > 147.83.194.21.80: . ack 1719 win 77 <nop,nop,timestamp
531419156 3872304345>
09:33:02.560256 IP 147.83.194.21.80 > 147.83.34.125.24374: . 1719:3167(1448) ack 627 win 1761
<nop,nop,timestamp 3872304345 531419156>
09:33:02.560261 IP 147.83.34.125.24374 > 147.83.194.21.80: . ack 3167 win 100 <nop,nop,timestamp
531419156 3872304345>
...

```

## Unit 3. Point to Point Protocols -TCP

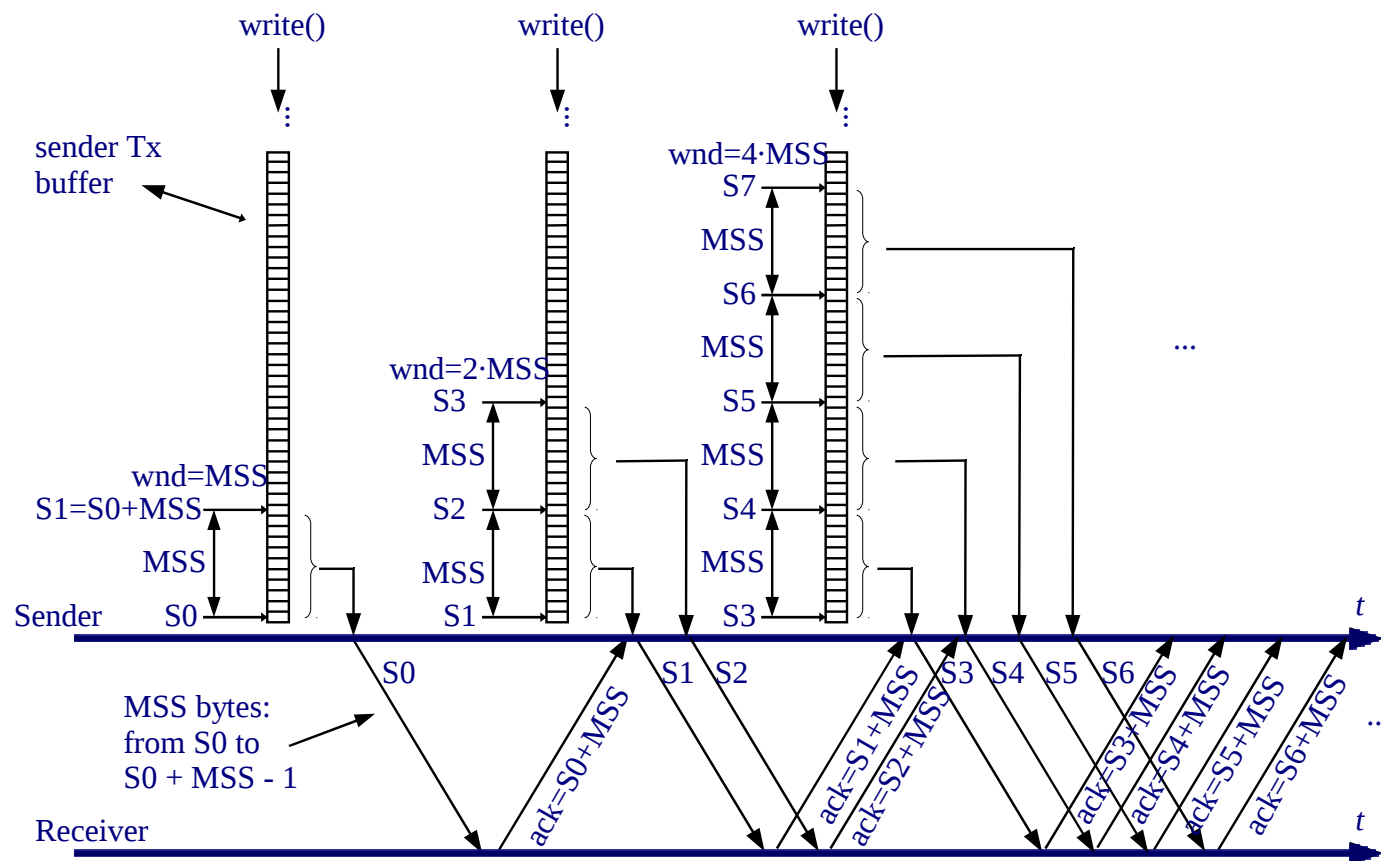
### TCP Protocol – TCP Options

- **Maximum Segment Size (MSS)**: Used in the TWH to initialize the MSS.
- **Window Scale factor**: Used in the TWH. The awnd is multiplied by  $2^{\text{Window Scale}}$  (i.e. the window scale indicates the number of bits to left-shift awnd). It allows using awnd larger than  $2^{16}$  bytes.
- **Timestamp**: Used to compute the Round Trip Time (RTT). Is a 10 bytes option, with the timestamp clock of the TCP sender, and an echo of the timestamp of the TCP segment being ack.
- **SACK**: In case of errors, indicate blocks of consecutive correctly received segments for Selective ReTx.

# Unit 3. Point to Point Protocols -TCP

## TCP Protocol – TCP Sequence Numbers

- The **sequence number** identifies the first payload byte.
- The **ack number** identifies the next byte the receiver is waiting for.



# Unit 3. Point to Point Protocols -TCP

## TCP Protocol – C Code example

```
main(int argc, char *argv[])
{
    int sock ;
    struct sockaddr_in serv_addr ;
    struct hostent *host ;

    if(argc != 2) {
        fprintf(stderr, "usage: %s hostname\n", argv[0]) ;
        exit(1) ;
    }
    host = gethostbyname(argv[1]) ; /* call the resolver for server addr. */
    if(host == NULL) {
        perror("gethostbyname ") ;
        exit(2) ;
    }
    bzero(&serv_addr, sizeof(serv_addr)) ;
    serv_addr.sin_family = AF_INET ;
    memcpy(&serv_addr.sin_addr, host->h_addr, host->h_length) ;
    serv_addr.sin_port = htons(3333) ;

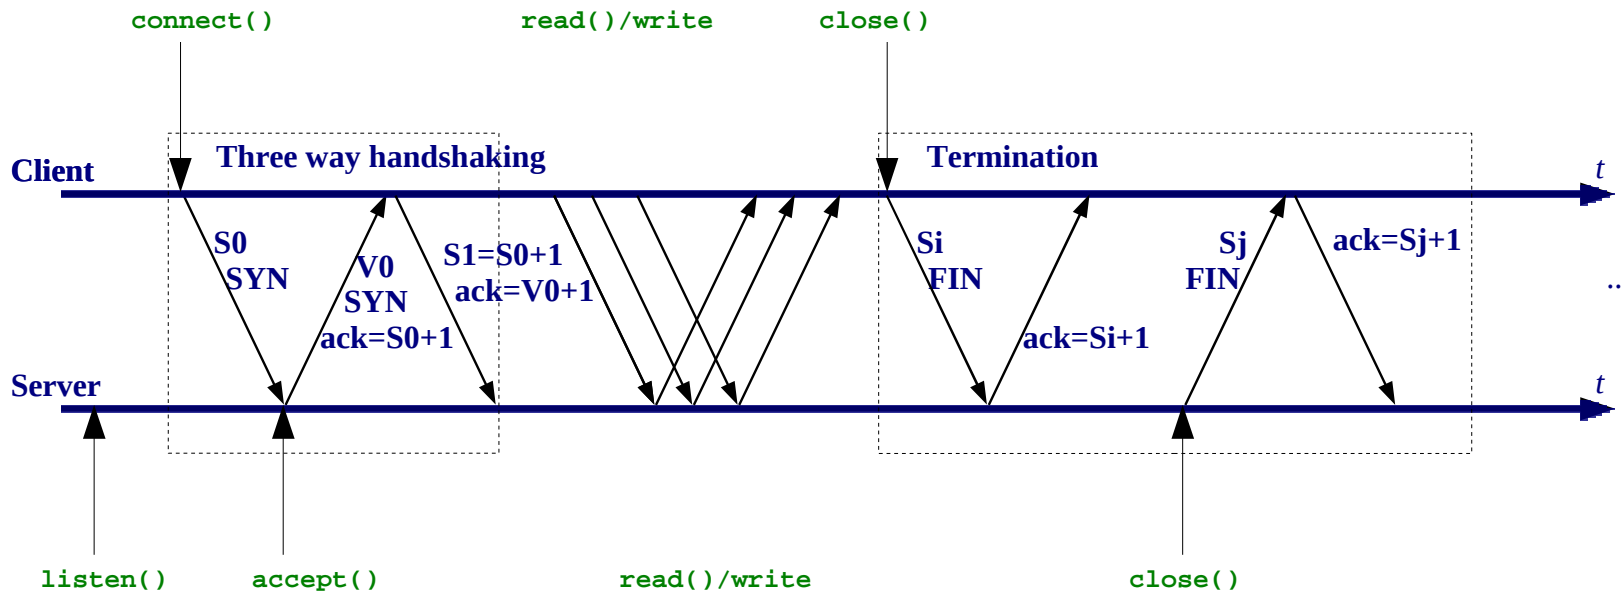
    sock = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP) ; /* create a socket */
    connect(sock, (struct sockaddr *)&serv_addr, sizeof(serv_addr)) ; /* initiate the connection */
    char msg[] = "hello world\n" ;
    write(sock, msg, strlen(msg)) ; /* write to TCP socket */
    char buf ;
    while(read(sock, &buf, 1) > 0) { /* read from TCP socket */
        printf("%c", buf) ;
    }
    close(sock) ; /* close the socket */
}
```



## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Connection Setup and Termination

- The **client** always send the 1<sup>st</sup> segment.
- **Three-way handshaking** segments have payload = 0.
- SYN and FIN segments **consume a sequence number**.
- Initial sequence number is **random**.



# Unit 3. Point to Point Protocols -TCP

## TCP Protocol – tcpdump example (web page download)

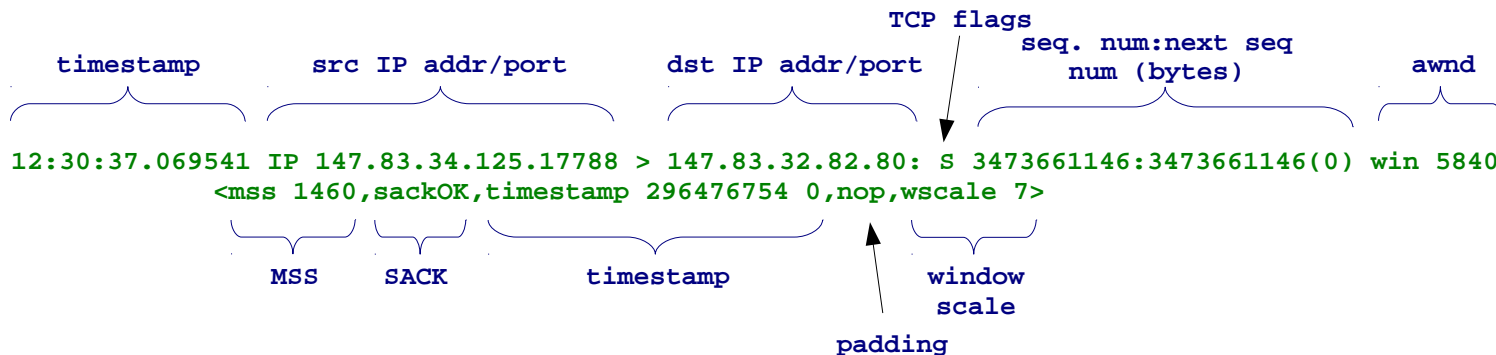
**TWH**

```

12:30:37.069541 IP 147.83.34.125.17788 > 147.83.32.82.80: S 3473661146:3473661146(0) win 5840 <mss
1460,sackOK,timestamp 296476754 0,nop,wscale 7>
12:30:37.070021 IP 147.83.32.82.80 > 147.83.34.125.17788: S 544373216:544373216(0) ack 3473661147 win 5792 <mss
1460,sackOK,timestamp 1824770623 296476754,nop,wscale 2>
12:30:37.070038 IP 147.83.34.125.17788 > 147.83.32.82.80: . ack 1 win 46 <nop,nop,timestamp 296476754
1824770623>
12:30:37.072763 IP 147.83.34.125.17788 > 147.83.32.82.80: P 1:602(601) ack 1 win 46 <nop,nop,timestamp 296476754
1824770623>
12:30:37.073546 IP 147.83.32.82.80 > 147.83.34.125.17788: . ack 602 win 1749 <nop,nop,timestamp 1824770627
296476754>
12:30:37.075932 IP 147.83.32.82.80 > 147.83.34.125.17788: P 1:526(525) ack 602 win 1749 <nop,nop,timestamp
1824770629 296476754>
12:30:37.075948 IP 147.83.34.125.17788 > 147.83.32.82.80: . ack 526 win 54 <nop,nop,timestamp 296476755
1824770629>
12:30:53.880704 IP 147.83.32.82.80 > 147.83.34.125.17788: F 526:526(0) ack 602 win 1749 <nop,nop,timestamp
1824787435 296476755>
12:30:53.920354 IP 147.83.34.125.17788 > 147.83.32.82.80: . ack 527 win 54 <nop,nop,timestamp 296480966
1824787435>
12:30:56.070200 IP 147.83.34.125.17788 > 147.83.32.82.80: F 602:602(0) ack 527 win 54 <nop,nop,timestamp
296481504 1824787435>
12:30:56.070486 IP 147.83.32.82.80 > 147.83.34.125.17788: . ack 603 win 1749 <nop,nop,timestamp 1824789625
296481504>


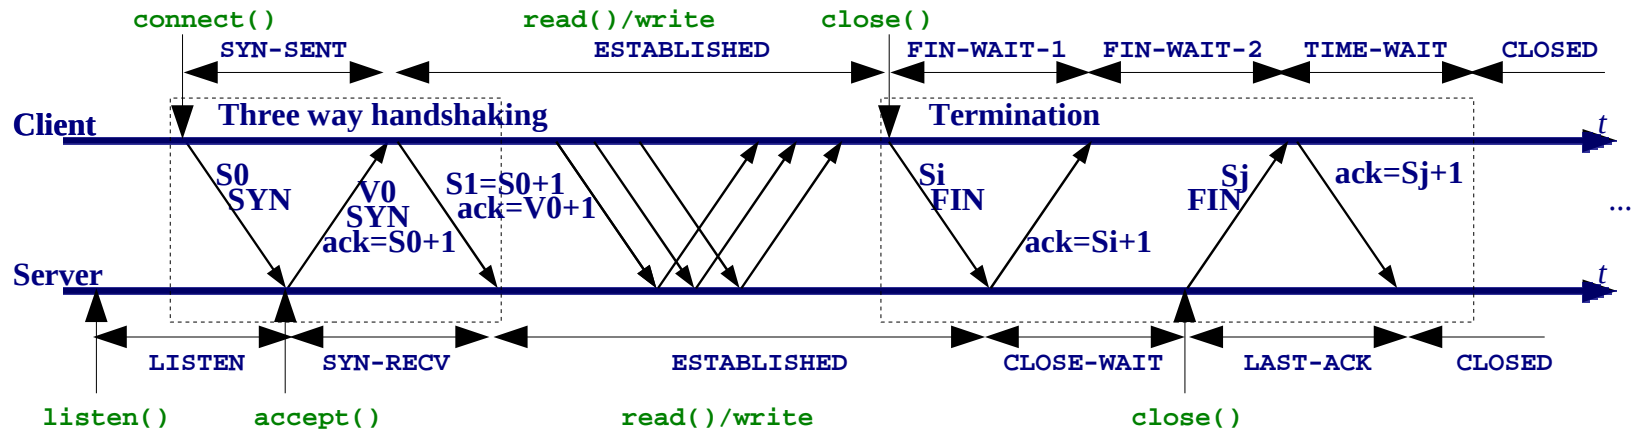
```

**Termination**



# Unit 3. Point to Point Protocols -TCP

## TCP Protocol – State diagram (simplified)



# Unit 3. Point to Point Protocols -TCP

## TCP Protocol – netstat dump

- Option -t shows tcp sockets.

```
linux# netstat -nt
Active Internet connections (w/o servers)
```

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State
tcp	0	1286	192.168.0.128:29537	199.181.77.52:80	ESTABLISHED
tcp	0	0	192.168.0.128:13690	67.19.9.2:80	TIME_WAIT
tcp	0	1	192.168.0.128:12339	64.154.80.132:80	FIN_WAIT1
tcp	0	1	192.168.0.128:29529	199.181.77.52:80	SYN_SENT
tcp	1	0	192.168.0.128:17722	66.98.194.91:80	CLOSE_WAIT
tcp	0	0	192.168.0.128:14875	210.201.136.36:80	ESTABLISHED
tcp	0	0	192.168.0.128:12804	67.18.114.62:80	ESTABLISHED
tcp	0	1	192.168.0.128:25232	66.150.87.2:80	LAST_ACK
tcp	0	0	192.168.0.128:29820	66.102.9.147:80	ESTABLISHED
tcp	0	0	192.168.0.128:29821	66.102.9.147:80	ESTABLISHED
tcp	1	0	127.0.0.1:25911	127.0.0.1:80	CLOSE_WAIT
tcp	0	0	127.0.0.1:25912	127.0.0.1:80	ESTABLISHED
tcp	0	0	127.0.0.1:80	127.0.0.1:25911	FIN_WAIT2
tcp	0	0	127.0.0.1:80	127.0.0.1:25912	ESTABLISHED

man netstat


The count of bytes not acknowledged by the remote host.

The count of bytes not copied by the user program connected to this socket.

## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Congestion Control (RFC 2581)

- $\text{window} = \min(\text{awnd}, \text{cwnd})$ 
  - The advertised window (**awnd**) is used for **flow control**.
  - The congestion window (**cwnd**) is used for **congestion control**.
- TCP interprets losses as congestion:



- Basic **Congestion Control** Algorithms:
  - Slow Start / Congestion Avoidance (**SS/CA**)
  - Fast Retransmit / Fast Recovery (**FR/FR**)

# Unit 3. Point to Point Protocols -TCP

## TCP Protocol – Slow Start / Congestion Avoidance (SS/CA)

- Variables:
  - `snd_una`: First non ack segment.
  - `ssthresh`: Threshold between SS and CA.

Initialization:

```
cwnd = MSS ; NOTE: RFC 2581 allows an initial window of 2 segments.
ssthresh = infinity ;
```


Each time an **ack confirming new data** is received:

```
if(cwnd < ssthresh) {
    cwnd += MSS ; /* Slow Start */
} else {
    cwnd += MSS * MSS / cwnd ; /* Congestion Avoidance */
}
```

When there is a **time-out**:

```
Retransmit snd_una ;
cwnd = MSS ;
ssthresh = max(min(ssthresh, cwnd) / 2, 2 MSS) ;
```

Time-out Example:



## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Slow Start / Congestion Avoidance (SS/CA)

- During **SS** cwnd is rapidly increased to the “operational point”.
- During **CA** cwnd is slowly increased looking for more available bandwidth.

Initialization:


```
cwnd = MSS ;
ssthresh = infinit ;
```

Each time an ack confirming new data is received:

```
if(cwnd < ssthresh) {
    cwnd += MSS ; /* SS */
} else {
    cwnd += MSS * MSS / cwnd ; /* CA */
}
```

When there is a time-out:

```
Retransmit snd_una ;
cwnd = MSS ;
ssthresh = max(min(awnd, cwnd) / 2, 2 MSS) ;
```



## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Fast Retransmit / Fast Recovery (FR/FR)

- Several algorithms are typically used to improve TCP performance.
- A TCP implementation with SS/CA, FR/FR is referred to as **TCP-Reno**.
- Other improvements include e.g. **SACK**.
- **FR/FR** basis: Duplicate acks are indication of losses.

Each ack arrival:

```

if(it is a duplicated ack) {
    if(it is the 3th duplicated ack) {
        retransmit snd_una ;
        ssthresh = max(min(awnd, cwnd) / 2, 2 MSS) ;
        cwnd = ssthresh + 3 MSS ;
        fast_recovery = TRUE ;
    } else if(fast_recovery == TRUE) {
        cwnd += MSS ;
    }
} else { /* if new data is ack */
    if(fast_recovery == TRUE) {
        cwnd = ssthresh ;
        fast_recovery = FALSE ;
    } else {
        /* Slow Start / Congestion Avoidance */
    }
}

```

Annotations:


- 3th dup-ack arrival (points to the 3th duplicated ack condition)
- cwnd is set to the wnd estimated previous to the loss + 3 (for the dup acks) (points to `cwnd = ssthresh + 3 MSS ;`)
- allows sending a new segment for each segment arriving to the TCP receiver (points to the `fast_recovery == TRUE` condition in the else branch)
- exit FR when new data is ack. (points to the `fast_recovery = FALSE ;` line)



# Unit 3. Point to Point Protocols -TCP

## TCP Protocol – Fast Retransmit / Fast Recovery (FR/FR)

Time axis at the sender side



Each ack arrival:

```

if(it is a duplicated ack) {
    if(it is the 3th duplicated ack) {
        retransmit snd_una ;
        ssthresh = max(min(awnd, cwnd) / 2, 2 MSS) ;
        cwnd = ssthresh + 3 MSS ;
        fast_recovery = TRUE ;
    } else if(fast_recovery == TRUE) {
        cwnd += MSS ;
    }
} else { /* if new data is ack */
    if(fast_recovery == TRUE) {
        cwnd = ssthresh ;
        fast_recovery = FALSE ;
    } else {
        /* Slow Start / Congestion Avoidance */
    }
}

```

## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Retransmission time-out (RTO)

- Activation:
  - RTO is active whenever there are pending acks.
  - When RTO is active, it is **continuously decreased**, and a ReTx occurs when RTO reaches zero.
  - Each time an **ack confirming new data** arrives:
    - RTO is computed.
    - RTO is restarted if there are pending acks, otherwise, RTO is stopped.
- Computation:
  - The TCP sender measures the RTT mean (**srtt**) and variance (**rttvar**).
  - The retransmission time-out is given by:  **$RTO = srtt + 4 \times rttvar$** .
  - RTO is **duplicated each retransmitted segment** (exponential backoff).
- RTT measurements:
  - Using “**slow-timer** tics” (coarse).
  - Using the TCP **timestamp option**.

## Unit 3. Point to Point Protocols -TCP

### TCP Protocol – Retransmission time-out (RTO)

